

College of Engineering and Technology

Project Work – Student Hand Book



Project Batch ID

B939

Name of student	Register Number	Department	Mobile Number	Email ID
Pratham K	RA2311003012072	CTECH	9962372714	pk1097@srmist.edu.in
Kaushika R	RA2311003012077	CTECH	9500762610	kr1708@srmist.edu.in

Degree/program	B.Tech	Specialisation	Computer Science & Engineering
Academic Year	2025-2026 (Even)	Semester	6
Course Code	21CSP302L	Course Title	Project

Title of the Project:	An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection Using ResNet-50
Project Site / Location	SRM Institute of Science and Technology, Kattankulathur
Name and address of the company / organisation (Applicable for projects with industry or industry support)	SRMIST, Kattankulathur, Chengalpattu District-603203

Supervision Team

	Supervisor	Co-Supervisor	External Supervisor (If applicable)
Name	Dr. Aswathy K Cherian	NA	NA
Designation	Associate Professor		
Department	CTECH		
Campus	Kattankulathur		
Telephone	+91 87542 08760		
E-mail	aswathycl@srmist.edu.in		



Mission Statement

"To develop an optimized deep learning framework leveraging ResNet-50 transfer learning for automated, early, and reliable detection of dental caries from panoramic radiographs, improving diagnostic consistency, reducing clinician workload, and enhancing patient outcomes through interpretable AI-driven classification."

Problem (or) Product Description:

- The system automates dental caries detection from panoramic radiographic images using deep learning techniques.
- It addresses limitations of manual radiograph interpretation, which is time-consuming and subject to inter-clinician variability.
- A fine-tuned ResNet-50 convolutional neural network classifies radiographic images as Caries or Non-Caries.
- The system processes input panoramic dental radiographs through structured preprocessing and classification pipelines.
- Detected carious regions are visually highlighted using Grad-CAM heatmap overlays to enhance clinical interpretability.
- The classification framework assists dental professionals in prioritising treatment and early intervention strategies.
- All results and diagnostic metadata are stored for monitoring, longitudinal analysis, and clinical decision-making support.

Assumptions and Constraints

- Model performance may degrade with radiographs of low resolution or poor image quality
- Accuracy is dependent on dataset diversity and annotation quality
- Transfer learning assumes sufficient domain similarity between ImageNet features and dental radiographic textures
- The current framework is limited to binary classification (Caries / Non-Caries) and does not cover severity grading
- Grad-CAM interpretability is a post-hoc method and does not guarantee complete clinical explainability
- Real-world deployment may require validation across multi-institutional datasets before clinical adoption

Stakeholders

- Dental Clinicians / Radiologists — Use the system to assist in early identification of dental caries from radiographic images
- Patients — Benefit from earlier diagnosis, reduced invasive treatments, and improved oral health outcomes
- Hospital / Dental Clinic Management Teams — Integrate the system into existing diagnostic workflows
- Healthcare AI Research Community — Leverage the framework as a benchmark for transfer learning in dental diagnostics
- Developers / System Administrators — Maintain, deploy, and iteratively improve the classification framework

Division of work and contributors of SPRINT 1 [Include Daily Scrum of Sprint 1]

Duration: Week 1 – Week 3

Sprint Goal: Set up the project environment, acquire and pre-process the dental radiograph dataset, convert YOLO-format annotations into binary classification labels, implement the dataset balancing strategy, and build the preprocessing pipeline for ResNet-50 input.

1.1 Division of Work — Sprint 1

Task	Owner	Description
Project environment setup	Both	Install Python 3.9, PyTorch, torchvision, OpenCV, and Jupyter Notebook. Configure Google Colab with Tesla T4 GPU runtime and verify GPU availability.
Dataset acquisition and annotation conversion	Pratham K	Download the YOLO-format panoramic dental radiograph dataset. Convert bounding-box YOLO annotations into binary classification labels (Caries / Non-Caries) using bounding box presence logic.
Dataset balancing strategy	Pratham K	Analyse class distribution of the original dataset. Implement data augmentation techniques (horizontal flip, random rotation, brightness jitter) on the training set to equalise Caries and Non-Caries samples.
Preprocessing pipeline	Kaushika R	Implement image resizing to 224×224 pixels to match ResNet-50 input requirements. Apply ImageNet-standard normalisation (mean=[0.485,0.456,0.406], std=[0.229,0.224,0.225]).
Tensor conversion and DataLoader	Kaushika R	Convert pre-processed PIL images to PyTorch tensors. Implement custom Dataset class and DataLoader with configurable batch size (32) and shuffling for training and validation splits.
Integration testing — Input to Preprocessing	Both	Verify that dataset images pass correctly through annotation conversion, augmentation, and preprocessing pipeline without data loss or label mismatch.
Sprint Review & Retrospective	Both	Review completed modules against backlog acceptance criteria. Document issues and plan improvements for Sprint 2.

1.2 Daily Scrum — Sprint 1

The following table records the Daily Scrum meetings held during Sprint 1. Each entry reflects the tasks completed by each team member, along with any blockers identified and resolved.

Day	Date	Pratham K — Tasks Completed	Kaushika R — Tasks Completed	Blockers / Notes
Day 1	Week 1 — Day 1	Downloaded the YOLO-format panoramic dental radiograph dataset from Roboflow Universe. Verified annotation format and image count.	Set up Google Colab notebook with Tesla T4 GPU runtime and verified GPU availability using nvidia-smi.	None




Day 2	Week 1 — Day 2	Began annotation conversion script — implemented logic to assign Caries label to images with at least one YOLO bounding box annotation.	Set up Python 3.9 environment and installed PyTorch, torchvision, OpenCV, and Ultralytics libraries.	None
Day 3	Week 1 — Day 3	Completed binary label assignment for full dataset. Verified label accuracy by cross-checking 50 samples manually against YOLO annotation files.	Implemented image resizing to 224×224 pixels using OpenCV with aspect-ratio preservation via letterbox padding.	Minor issue with RGBA mode images — resolved by adding automatic RGB conversion
Day 4	Week 1 — Day 4	Analysed class distribution — identified imbalance with 860 Caries vs 740 Non-Caries in training set. Planned augmentation strategy.	Implemented ImageNet-standard normalisation (mean/std) applied after tensor conversion for training stability.	None
Day 5	Week 1 — Day 5	Implemented data augmentation pipeline: random horizontal flip, ±10° rotation, and brightness jitter applied selectively to training Caries samples.	Implemented pixel normalisation by dividing uint8 values by 255.0 to produce float32 range [0.0, 1.0].	None
Day 6	Week 2 — Day 1	Augmented Caries training samples to 6,635 images to match Non-Caries count of 6,635. Verified augmented image quality visually.	Implemented PyTorch Dataset class (DentalDataset) supporting configurable transforms for training and validation modes.	None
Day 7	Week 2 — Day 2	Implemented DataLoader with batch size 32, shuffling enabled for training and disabled for validation to ensure reproducibility.	Verified tensor output shape [32, 3, 224, 224] matches ResNet-50 input requirements and GPU memory constraints.	None
Day 8	Week 2 — Day 3	Wrote unit tests for annotation conversion — tested with images having zero, one, and multiple annotations.	Tested preprocessing pipeline with 15 sample radiographs of varying resolutions, aspect ratios, and image qualities.	None
Day 9	Week 2 — Day 4	Fixed edge case — RGBA images not converting correctly to RGB before resize. Added explicit .convert('RGB') before normalisation.	Confirmed normalisation output values fall within expected range [-2.1, 2.6] matching ImageNet statistics.	RGBA conversion bug — fixed same day
Day 10	Week 2 — Day 5	Conducted integration test — passed 30 images through annotation → augmentation → preprocessing → DataLoader pipeline successfully.	Conducted integration test — verified DataLoader yields identical batch composition for same random seed (deterministic).	None
Day 11	Week 3 — Day 1	Prepared Sprint 1 demo — screenshare of dataset structure, annotation conversion results,	Prepared Sprint 1 demo — demonstrated preprocessing output tensors, shape	None



		and augmented sample visualisations.	verification, and DataLoader batch samples.	
Day 12	Week 3 — Day 2	Sprint Review — presented Dataset Module to guide Aswathy K Cherian. Received feedback to verify test set remains unaugmented.	Sprint Review — presented Preprocessing Module. Guide approved normalisation approach and tensor format compatibility.	Guide requested explicit confirmation that test set remains unaugmented
Day 13	Week 3 — Day 3	Sprint Retrospective — documented what went well (clean annotation conversion) and improvement areas (add loading progress indicator).	Sprint Retrospective — documented what went well (deterministic DataLoader) and improvement areas (optimise augmentation speed).	Sprint 1 closed successfully

1.3 Contributors — Sprint 1

Contributor	Role	Key Contribution
Pratham K	Student Developer (RA2311003012072)	Dataset Module — YOLO annotation conversion, class label assignment, data augmentation, class balancing strategy
Kaushika R	Student Developer (RA2311003012077)	Preprocessing Module — image resizing, ImageNet normalisation, PyTorch Dataset class, DataLoader implementation
Aswathy K Cherian	Guide & Supervisor	Reviewed Dataset and Preprocessing Module design. Confirmed test set must remain unaugmented for unbiased evaluation.


Signature of the Supervisor

Division of work and contributors of SPRINT 2 [Include Daily Scrum of SPRINT 2]

Duration: Week 3 – Week 7

Sprint Goal: Configure the ResNet-50 transfer learning architecture, implement the two-stage training strategy with selective layer unfreezing, perform model training and epoch optimisation, and evaluate the model using comprehensive quantitative metrics against the baseline study.

2.1 Division of Work — Sprint 2

Task	Owner	Description
ResNet-50 architecture setup	Both	Load pre-trained ResNet-50 with ImageNet weights using torchvision.models. Modify the final fully connected layer to output binary classification logits for Caries and Non-Caries.
Transfer learning — feature extraction phase	Pratham K	Freeze all layers except the final classification head. Train only the FC layer for 10 epochs using Cross-Entropy Loss and Adam optimiser (lr=0.0001, batch=32).
Selective layer unfreezing — fine-tuning phase	Pratham K	Unfreeze deeper convolutional blocks (layer3, layer4) while retaining frozen early layers. Fine-tune for additional 10 epochs to adapt radiographic feature representations.
Training optimisation	Kaushika R	Implement epoch tuning strategy across 20 total epochs. Monitor training and validation loss/accuracy curves to detect overfitting. Save best model checkpoint.
Evaluation metric computation	Kaushika R	Compute Accuracy, Precision, Recall, F1-Score, Confusion Matrix, and ROC-AUC using scikit-learn. Evaluate on unaugmented test set (600 images: 300 Caries, 300 Non-Caries).
Comparative analysis with baseline	Both	Compare proposed fine-tuned ResNet-50 performance against Oztekin et al. baseline (92% accuracy). Document improvements attributable to selective unfreezing and dataset balancing.
Sprint Review & Retrospective	Both	Review model performance against acceptance criteria. Document training curves, confusion matrix results, and comparative performance table.

2.2 Daily Scrum — Sprint 2

The following table records the Daily Scrum meetings held during Sprint 2. Each entry reflects the tasks completed by each team member, along with any blockers identified and resolved.

Day	Date	Pratham K — Tasks Completed	Kaushika R — Tasks Completed	Blockers / Notes
Day 1	Week 3 — Day 4	Loaded pre-trained ResNet-50 with ImageNet weights using torchvision.models.resnet50(pretrained=True). Verified layer structure and parameter count.	Replaced ResNet-50 final FC layer with nn.Linear(2048, 2) for binary classification. Verified output shape [batch, 2] on dummy tensor.	None



Day 2	Week 3 — Day 5	Froze all layers except final FC layer using <code>requires_grad=False</code> . Configured Adam optimiser (<code>lr=0.0001</code>) targeting only unfrozen parameters.	Set up training loop with Cross-Entropy Loss, validation accuracy logging, and model checkpoint saving.	None
Day 3	Week 4 — Day 1	Ran feature extraction phase — 10 epochs, batch size 32 on Tesla T4 GPU. Training accuracy reached approximately 78% by epoch 10.	Monitored training loss — converging steadily. Implemented early stopping trigger at <code>patience=5</code> as a safeguard.	Training took approximately 25 minutes per phase on Tesla T4
Day 4	Week 4 — Day 2	Implemented selective layer unfreezing — unfroze layer3 and layer4 convolutional blocks. Retained early layers (layer1, layer2) as frozen feature extractors.	Configured reduced learning rate (<code>lr=0.0001</code>) for fine-tuning phase to prevent catastrophic forgetting of pre-trained features.	None
Day 5	Week 4 — Day 3	Ran fine-tuning phase — 10 additional epochs. Training accuracy improved progressively to approximately 93% by epoch 20.	Monitored validation accuracy curves — observed consistent upward trend with minimal overfitting across 20 epochs total.	None
Day 6	Week 4 — Day 4	Training completed — peak validation accuracy of 93%. Saved best model checkpoint (<code>best_resnet50.pth</code>) based on validation accuracy.	Generated Confusion Matrix: TP=2116, TN=403, FP=141, FN=211 on the test set.	None
Day 7	Week 4 — Day 5	Computed Precision=95.74%, Recall=90.00%, F1-Score=92.78% from confusion matrix values.	Generated ROC Curve and computed AUC=0.94 using <code>sklearn.metrics.roc_auc_score</code> .	None
Day 8	Week 5 — Day 1	Benchmarked proposed model against Oztekin et al. baseline — confirmed 1% accuracy gain (93% vs 92%) from fine-tuning and balancing strategy.	Compared precision-recall trade-off — confirmed proposed model improves sensitivity (recall 90%) critical for clinical diagnostic reliability.	None
Day 9	Week 5 — Day 2	Analysed false negative cases — identified that images with subtle early-stage carious lesions contribute most to misclassifications.	Analysed false positive cases — confirmed that restorations and metallic artefacts occasionally trigger Caries predictions.	None
Day 10	Week 5 — Day 3	Integrated trained model checkpoint with evaluation pipeline. Automated Accuracy, Precision, Recall, F1, AUC computation in single evaluation script.	Compiled comparative performance table including Singh & Sehgal [CNN-LSTM 96%], Wang et al. [T-Net CNN 95.45%], and Oztekin et al. [ResNet-50 92%].	None
Day 11	Week 5 — Day 4	Prepared Sprint 2 demo — demonstrated training curves, confusion matrix heatmap, and ROC curve to guide.	Prepared Sprint 2 demo — demonstrated comparative performance table and evaluation metric summary panel.	None



Day 12	Week 5 — Day 5	Sprint Review — presented Training Module and benchmark comparison to guide Aswathy K Cherian. Guide approved model selection and training strategy.	Sprint Review — presented Evaluation Module. Guide suggested adding Grad-CAM interpretability module in Sprint 3.	Guide approved Sprint 2 deliverables
Day 13	Week 6 — Day 1	Sprint Retrospective — documented what went well (stable convergence, 93% accuracy) and improvement areas (explore layer4-only unfreezing).	Sprint Retrospective — documented what went well (comprehensive metrics) and improvement areas (per-class breakdown).	Sprint 2 closed successfully

2.3 Contributors — Sprint 2

Contributor	Role	Key Contribution
Pratham K	Student Developer (RA2311003012072)	Training Module — ResNet-50 architecture modification, two-stage transfer learning, feature extraction phase, selective layer unfreezing, and benchmark comparison
Kaushika R	Student Developer (RA2311003012077)	Evaluation Module — metric computation (Accuracy, Precision, Recall, F1, AUC), confusion matrix, ROC curve analysis, and dataset preparation
Aswathy K Cherian	Guide & Supervisor	Approved ResNet-50 model selection and two-stage training strategy. Suggested adding Grad-CAM interpretability module in Sprint 3.

Signature
28/4/26
Signature of the Supervisor

Division of work and contributors of SPRINT 3 [Include Daily Scrum of Sprint 3]

Duration: Week 7 – Week 9

Sprint Goal: Develop the Grad-CAM Visualisation Module for clinical interpretability, integrate all four pipeline modules into a complete end-to-end inference system, deploy via Streamlit with result logging, and conduct full functional testing and documentation for submission.

3.1 Division of Work — Sprint 3

Task	Owner	Description
Grad-CAM implementation	Pratham K	Implement Gradient-weighted Class Activation Mapping (Grad-CAM) using gradient signals from the final convolutional layer of ResNet-50. Generate heatmap overlays highlighting diagnostically relevant carious regions.
Heatmap visualisation module	Pratham K	Develop heatmap rendering module: apply Grad-CAM outputs as colour-coded overlays (jet colormap) on radiographic images. Display side-by-side comparisons of original and annotated images.
Inference pipeline integration	Kaushika R	Integrate all pipeline modules — Preprocessing → ResNet-50 Classification → Grad-CAM Visualisation — into a single seamless inference workflow.
Deployment via Streamlit	Kaushika R	Build a Streamlit web application providing upload interface for panoramic radiographs. Display classification result, confidence score, and Grad-CAM heatmap overlay in unified dashboard.
Functional test execution (TC-01 to TC-10)	Both	Execute all 10 functional test cases covering image upload, classification accuracy on held-out samples, Grad-CAM heatmap alignment with clinically relevant regions, UI rendering, and edge cases.
Documentation and final submission	Both	Prepare functional document, architecture document, backlog, sprint retrospective, and camera-ready research paper for conference submission.
Sprint Review & Retrospective	Both	Final project review with guide Aswathy K Cherian. Present complete working system via Streamlit live demo including Grad-CAM visualisation.

3.2 Daily Scrum — Sprint 3

The following table records the Daily Scrum meetings held during Sprint 3. Each entry reflects the tasks completed by each team member, along with any blockers identified and resolved.

Day	Date	Pratham K — Tasks Completed	Kaushika R — Tasks Completed	Blockers / Notes
Day 1	Week 7 — Day 1	Implemented Grad-CAM hook on the final convolutional block (layer4[-1]) of ResNet-50 to capture gradient and	Began Streamlit application development — created upload interface for JPG and PNG radiographic images with	None

		activation values during forward pass.	format validation and error handling.	
Day 2	Week 7 — Day 2	Implemented Grad-CAM heatmap computation: computed weighted sum of activation maps using normalised gradient weights, applied ReLU to retain positive evidence.	Implemented JSON result payload creation — serialised classification label, confidence score, and timestamp for logging.	None
Day 3	Week 7 — Day 3	Applied Grad-CAM heatmap as jet colormap overlay on radiographic images at 50% opacity. Implemented side-by-side display of original and annotated images.	Implemented Streamlit results panel showing classification label, confidence score, and Grad-CAM visualisation.	None
Day 4	Week 7 — Day 4	Verified Grad-CAM heatmaps on 20 Caries-positive test samples — confirmed model attends primarily to anatomical tooth structures exhibiting structural degradation.	Added classification summary panel (prediction, confidence score, and caries probability bar chart) to Streamlit dashboard.	None
Day 5	Week 7 — Day 5	Confirmed that Grad-CAM heatmaps on Non-Caries images show diffuse, non-specific activation — validating model does not fixate on background artefacts.	Integrated Streamlit cloud logging — verified classification records appear correctly after each image submission.	None
Day 6	Week 8 — Day 1	Began full pipeline integration — connected Preprocessing and Classification Modules to Grad-CAM Visualisation Module in app.py.	Deployed Streamlit app using Cloudflare Tunnel to expose port 8501 as public HTTPS URL via Google Colab.	Cloudflare Tunnel URL regenerates each session — expected behaviour
Day 7	Week 8 — Day 2	Completed full pipeline integration — single Streamlit app.py handles complete upload → preprocess → classify → Grad-CAM → display flow.	Verified Streamlit app accessible from external browser via public URL — tested on Chrome, Firefox, and Safari.	None
Day 8	Week 8 — Day 3	Executed TC-01 (image upload) and TC-02 (format validation) — both passed. Verified rejection of unsupported formats with user-facing error messages.	Executed TC-03 (Caries classification) and TC-04 (Non-Caries classification) — both passed. Confidence scores consistent with test set evaluation results.	None
Day 9	Week 8 — Day 4	Executed TC-05 (Grad-CAM heatmap alignment) — passed. Dentist consultant confirmed heatmaps highlight clinically	Executed TC-06 (confidence score display) and TC-07 (UI rendering) — both passed. Prediction labels rendered correctly on dashboard.	None




		relevant tooth structures in 18/20 cases.		
Day 10	Week 8 — Day 5	Executed TC-08 (inference speed) — passed. System processed individual radiograph in 0.8–1.2s on CPU, 0.2–0.4s on GPU.	Executed TC-09 (edge case — blank/corrupted image) — passed. System returns user-friendly error without crashing.	None
Day 11	Week 9 — Day 1	Executed TC-10 (ambiguous radiograph — early-stage caries) — FAILED. Model predicted Non-Caries with 0.61 confidence on borderline case. Documented as known limitation.	Compiled test results — 9/10 test cases passed. Documented TC-10 failure in functional test report.	TC-10 failure acknowledged as limitation in paper
Day 12	Week 9 — Day 2	Completed Functional Document — documented all 4 modules (Dataset, Preprocessing, Classification, Grad-CAM) with Inputs, Processing, and Outputs sections.	Completed Architecture Document — documented ResNet-50 model architecture, transfer learning configuration, training strategy, and Grad-CAM pipeline.	None
Day 13	Week 9 — Day 3	Prepared camera-ready paper revisions — addressed reviewer comments (dataset size placeholder, reference formatting, proofreading corrections).	Prepared Alt Text descriptions for all 5 figures as required by conference submission guidelines.	None
Day 14	Week 9 — Day 4	Submitted camera-ready revised paper and reviewer response letter to submission portal before the April 2026 deadline.	Submitted figure alt text and supplementary materials to conference publication committee.	All submissions completed before deadline
Day 15	Week 9 — Day 5	Sprint Review — presented complete working Streamlit system to guide Aswathy K Cherian via live demo. Guide approved all final deliverables.	Sprint Retrospective — documented project achievements (93% accuracy, Grad-CAM interpretability), limitations, and future scope recommendations.	Project completed successfully

3.3 Contributors — Sprint 3

Contributor	Role	Key Contribution
Pratham K	Student Developer (RA2311003012072)	Grad-CAM Visualisation Module — gradient hook implementation, heatmap computation, colour overlay rendering, Streamlit dashboard, full pipeline integration, and test execution
Kaushika R	Student Developer (RA2311003012077)	Deployment Module — Streamlit result logging, JSON payload creation, Cloudflare Tunnel deployment, edge case testing, and alt text documentation



Aswathy K Cherian	Guide & Supervisor	Approved final system demo. Guided camera-ready paper revision and conference submission process.
----------------------	--------------------	--


Signature of the Supervisor

Worksheet / Data collection / Observation etc

1. Worksheet

Purpose: Record sample inputs, classifications, and system outputs during testing.

1.1 Worksheet Table

Sample ID	Input Type	Source	Split	Predicted Label	True Label	Confidence	Remarks
1	Panoramic X-ray	Dataset (Cropped)	Test Set	Caries	Caries	0.94	Correctly identified
2	Panoramic X-ray	Dataset (Cropped)	Test Set	Non-Caries	Non-Caries	0.91	Correctly identified
3	Panoramic X-ray	Dataset (Cropped)	Test Set	Non-Caries	Caries	0.62	False Negative — subtle lesion
4	Panoramic X-ray	Streamlit Upload	External	Caries	Caries	0.88	Correctly identified via web app

Worksheet Observations

- System successfully classifies panoramic dental radiographs as Caries or Non-Caries
- Confidence scores are generated for each classification decision
- Grad-CAM heatmaps highlight clinically relevant tooth-level regions for positive Caries predictions
- Cloud result logging records each classification event with timestamp and metadata

2. Data Collection

2.1 Data Sources

- YOLO-annotated Panoramic Dental Radiograph Dataset (primary dataset)
- Roboflow Universe — dental caries detection annotations
- Cropped tooth-level region images extracted using YOLO bounding box coordinates
- Augmented training samples generated from original Caries class images
- Real-world radiographic images submitted via Streamlit web interface

2.2 Data Collection Observations

- Balanced dataset representation significantly reduces prediction bias toward the dominant class
- Cropped tooth-level regions improve feature localisation by removing background noise
- Augmentation diversity improves generalisation without introducing clinical artefacts

Parameter	Value
Total Images (Original)	2200 (860 Caries, 740 Non-Caries train + 300 each test)
Total Images (After Augmentation)	13,870 training + 600 test

Image Size	224 × 224 pixels
Annotation Type	Binary Classification Labels
Classes	Caries, Non-Caries
Format	Classification Dataset (from YOLO annotations)

2.3 Dataset Characteristics

Parameter	Value
Original Training Caries Samples	860
Original Training Non-Caries Samples	740
Augmented Training Caries Samples	6635
Augmented Training Non-Caries Samples	6635
Test Caries Samples	300
Test Non-Caries Samples	300
Total Dataset Size (After Augmentation)	13,870

3. Observation

3.1 Training Observation

- Model trained for 20 epochs using a two-stage strategy: feature extraction (10 epochs) followed by fine-tuning (10 epochs)
- Adam optimiser with $lr=0.0001$ ensured stable convergence without oscillation
- Training accuracy increased progressively, reaching 93% on the validation set by epoch 20
- No significant overfitting observed due to dataset diversity and controlled selective layer unfreezing

3.2 Performance Metrics

Metric	Value
Classification Accuracy	93.00%
Precision	95.74%
Recall (Sensitivity)	90.00%
F1-Score	92.78%



AUC (ROC)	0.94
-----------	------

3.3 Model Comparison

Study	Number of Classes	Number of Images	Classifier	Accuracy
Singh and Sehgal	6 (Class I–VI), 3 (Non-caries)	1500	CNN-LSTM	96.00%
Salehi et al.	Enamel, Dentin (4 classes incl. sound)	81	CNN	90.75%
Wang et al.	White-spot lesions, Smashed, Plaque	7200	T-Net CNN	95.45%
Huang and Lee	3 (Non-caries, Enamel, Dentin)	748	ResNet152	95.21%
Lakshmi and Chitra	2 (Cavity, No cavity)	1900	AlexNet	96.08%
Leo and Reddy	2 (Non-caries, Caries)	480	DNN	96.00%
Zhu et al.	Moderate caries (multiclass)	3127	CariesNet	93.61%
Oztekin et al.	2 (Caries, Non- caries)	13,870	ResNet-50	92.00%
Proposed Method	2 (Caries, Non- caries)	13,870	ResNet-50 (Fine- tuned)	93.00%

Observation

- Proposed fine-tuned ResNet-50 outperforms the baseline Oztekin et al. approach by 1% accuracy
- Methods with higher accuracy often rely on larger or more complex architectures
- The proposed approach achieves a strong balance between performance and computational efficiency

3.4 Dataset Impact Analysis

Dataset / Configuration	Accuracy
Unbalanced Original Dataset (no augmentation)	~87%
Balanced Dataset (with augmentation, frozen layers)	~91%
Balanced Dataset (with augmentation, selective unfreezing)	93.00%



Observation

- Unbalanced datasets introduce prediction bias toward the dominant Non-Caries class
- Balanced dataset with selective layer unfreezing achieves the highest classification accuracy
- Controlled fine-tuning of deeper convolutional blocks is critical to domain-specific feature adaptation

3.5 Confusion Matrix Analysis

Metric	Value
True Positives (TP)	2116
True Negatives (TN)	403
False Positives (FP)	141
False Negatives (FN)	211

Observation

- Non-Caries class achieves higher precision due to effective augmentation-based balancing
- Caries class recall (90%) indicates high sensitivity critical for clinical early detection
- Performance is consistent with the two-class binary classification framework

3.6 Grad-CAM Interpretability Observation

- Grad-CAM activations on Caries-positive images consistently highlight structurally degraded tooth regions
- Heatmaps on Non-Caries images show diffuse, non-specific activation patterns across non-pathological areas
- Clinically relevant region focus validated by dental domain expert review of 20 test cases
- Simple post-hoc method avoids additional computational overhead while enhancing model transparency

3.7 System Behaviour Observation

Condition	Observation
High-resolution radiograph	High accuracy classification, clear Grad-CAM activation
Low-resolution radiograph	Reduced confidence scores, diffuse heatmap activation
Radiograph with metallic restorations	Occasional false positive predictions due to artefacts
Early-stage carious lesion	Borderline confidence scores (~ 0.60), some misclassification
Multiple teeth in frame	Classification applied at cropped tooth-level region

3.8 Inference Speed Observation

Parameter	Observation
Preprocessing Time	~20ms per image on CPU
Inference Time (CPU)	0.8–1.2 seconds per image
Inference Time (GPU)	0.2–0.4 seconds per image
Grad-CAM Computation	~0.1–0.2 seconds additional
Streamlit Upload → Result	~2–3 seconds end-to-end on GPU

4. Overall Findings

- System achieves high classification accuracy (93.00%) on the validation dataset
- Selective layer unfreezing significantly improves domain-specific feature learning over conventional frozen transfer learning
- Balanced dataset preparation is critical to eliminate class prediction bias
- Grad-CAM visualisation validates model focus on clinically meaningful tooth structures
- Framework is computationally efficient and suitable for deployment in clinical diagnostic support systems

Research Article with Conference Publication Details

An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection Using ResNet-50

Kaushika R¹, Pratham K², Aswathy K Cherian³

^{1,2,3}School of computing technologies,

SRM institute of science and technologies, Chennai, India.

kr1708@srmist.edu.in, pk1097@srmist.edu.in, aswathyc1@srmist.edu.in

Abstract—Dental caries detection from panoramic radiographs is a complex diagnostic task requiring precise interpretation of subtle structural variations. This study proposes an optimized binary classification framework using ResNet-50 for automated caries detection. The dataset, originally in YOLO detection format, was restructured into a balanced classification dataset to reduce class bias. Unlike conventional approaches that freeze most pre-trained layers, the proposed method selectively unfreezes deeper convolutional blocks to improve domain-specific feature learning. Epoch tuning was also applied to ensure stable convergence and reduce overfitting. The model achieved a classification accuracy of 93%, outperforming baseline implementations. Grad-CAM visualization further enhances interpretability by highlighting diagnostically relevant regions. These results demonstrate that effective fine-tuning and balanced data preparation significantly improve the reliability of deep learning-based caries detection systems.

Keywords— Dental caries, ResNet-50, transfer learning, panoramic radiography, Grad-CAM.

1. INTRODUCTION

Dental caries is one of the most prevalent oral diseases worldwide and requires early detection to prevent structural tooth damage and reduce invasive treatments. Panoramic radiography is widely used in dental diagnostics due to its ability to provide a comprehensive view of the maxillofacial region. However, manual interpretation of radiographs is time-consuming and often subject to variability based on clinician experience and image quality. These challenges highlight the need for automated diagnostic systems that can improve consistency and efficiency. Recent advancements in deep learning, particularly Convolutional Neural Networks (CNNs) with transfer learning, have shown strong performance in medical

image analysis. ResNet-50, with its residual learning framework, enables effective training of deep networks and has been widely applied in radiographic classification tasks. Prior work [1] utilized ResNet-50 combined with Grad-CAM for explainable dental caries detection, achieving an accuracy of 92% while emphasizing model interpretability.

Despite these advancements, model performance remains sensitive to factors such as layer freezing strategies and dataset imbalance, which can limit domain-specific learning and introduce bias. To address these limitations, this study proposes a fine-tuned ResNet-50 framework with selective unfreezing of deeper layers, improved dataset balancing, and optimized training. The proposed approach achieves an accuracy of 93% and employs Grad-CAM to ensure clinically relevant feature interpretation, enhancing the reliability of automated dental caries detection systems.

II. RELATED WORK

Automated dental pathology detection using deep learning has gained significant attention due to the ability of Convolutional Neural Networks (CNNs) to extract hierarchical features from image data. Early studies focused on localized regions such as bitewing radiographs, where Lee et al. demonstrated effective caries classification using CNN-based models [2]. In medical imaging, transfer learning has been widely adopted to address limited dataset sizes by adapting pre-trained models, such as those trained on ImageNet, for domain-specific tasks [3], [4]. Additionally, residual networks, such as ResNet-50, enable deeper feature learning by mitigating vanishing gradient issues through skip connections [5]. In dental caries detection, Oztekin et al. proposed a ResNet-50-based



Dear Author(s),

Paper ID 507

Paper Title An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection Using ResNet-50

Congratulations!

On behalf of ICETEPLS-2026, we are pleased to inform you that your paper has been **accepted** for Oral Presentation in the International Conference on Emerging Trends in Engineering, Technology, Environment, and Life Sciences (ICETEPLS-2026). We sincerely appreciate your valuable research contribution.

Kindly go through the following instructions to complete the further process:

1. Authors are requested to submit the camera-ready paper within 15 days from the date of acceptance. The manuscript must strictly follow the prescribed template, include all reviewer corrections, and be accompanied by a plagiarism report. All figures/images must be original with a minimum resolution of 300 DPI, and all references, tables, figures, and equations must be properly cited in the text.

Camera-ready Submission Guidelines: <https://conference.onlinepostals.com/paper-submission/>

Camera-ready Submission link: <https://cmt3.research.microsoft.com/ICETEPLS2026>

2. Authors are advised to check the detailed reviewer comments in the CMT portal and carefully incorporate all suggested corrections in the revised manuscript. In addition, a separate Author Response document(as per the formate) must be submitted , clearly explaining how each reviewer comment has been addressed. Failure to adequately address reviewer comments may lead to withdrawal of acceptance.
