

# OPTIMIZED MACHINE LEARNING FRAMEWORK FOR DENTAL CARIES DETECTION

21CSP302L- PROJECT

*Submitted by*

KAUSHIKA R [RA2311003012077]  
PRATHAM K [RA2311003012072]

*Under the Guidance of*

Dr. ASWATHY K CHERIAN

(Associate Professor, Computing Technologies)

*in partial fulfillment of the requirements for the degree of*

BACHELOR OF TECHNOLOGY  
in  
COMPUTER SCIENCE AND ENGINEERING



DEPARTMENT OF COMPUTING TECHNOLOGIES SCHOOL OF  
ENGINEERING AND TECHNOLOGY SRM INSTITUTE OF SCIENCE  
AND TECHNOLOGY  
KATTANKULATHUR- 603 203

APRIL 2026



Department of Computing Technologies  
SRM Institute of Science &  
Technology Own Work  
Declaration Form

This sheet must be filled in (each box ticked to show that the condition has been met). It must be signed and dated along with your student registration number and included with all assignments you submit – work will not be marked unless this is done.

To be completed by the student for all assessments

Degree/ Course : B TECH COMPUTER SCIENCE AND ENGINEERING  
Student Name : KAUSHIKA R, PRATHAM K  
Registration Number : RA2311003012077, RA2311003012072  
Title of Work : OPTIMIZED MACHINE LEARNING FRAMEWORK FOR  
DENTAL CARIES DETECTION

We hereby certify that this assessment compiles with the University's Rules and Regulations relating to Academic misconduct and plagiarism, as listed in the University Website, Regulations, and the Education Committee guidelines.

We confirm that all the work contained in this assessment is my / our own except where indicated, and that We have met the following conditions:

- Clearly referenced / listed all sources as appropriate
- Referenced and put in inverted commas all quoted text (from books, web, etc)
- Given the sources of all pictures, data etc. that are not my own
- Not made any use of the report(s) or essay(s) of any other student(s) either past or present
- Acknowledged in appropriate places any help that I have received from others (e.g. fellow students, technicians, statisticians, external sources)
- Compiled with any other plagiarism criteria specified in the Course handbook /University website

I understand that any false claim for this work will be penalized in accordance with the University policies and regulations.

**DECLARATION:**

We are aware of and understand the University's policy on Academic misconduct and plagiarism and we certify that this assessment is my / our own work, except where indicated by referring, and that I have followed the good academic practices noted above.

*[Signature]* 22/4/26  
If you are working in a group, please write your registration numbers and sign with the date for every student in your group.


*[Signature]* 22/4/26




**SRM INSTITUTE OF SCIENCE AND TECHNOLOGY  
KATTANKULATHUR – 603 203**

**BONAFIDE CERTIFICATE**

Certified that 21CSP302L - Project report titled "**OPTIMIZED MACHINE LEARNING FRAMEWORK FOR DENTAL CARIES DETECTION**" is the bonafide work of "**KAUSHIKA R [RA2311003012077], PRATHAM K [RA2311003012072]**" who carried out the project work under my supervision. Certified further, that to the best of my knowledge the work reported herein does not form any other project report or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

  
28/4/26  
SIGNATURE

**Dr. ASWATHY K CHERIAN**  
**ASSOCIATE PROFESSOR**  
Department of Computing Technologies





**SIGNATURE**  
**Dr. NIRANJANA G**  
**PROFESSOR & HEAD**  
Department of Computing Technologies




  
28/4/26  
EXAMINER 1

Name & Signature:



  
28/4/26  
EXAMINER 2

Name & Signature:



## ACKNOWLEDGEMENTS

We express our humble gratitude to **Dr. C. Muthamizhchelvan**, Vice-Chancellor, SRM Institute of Science and Technology, for the facilities extended for the project work and his continued support.

We extend our sincere thanks to **Dr. Leenus Jesu Martin M**, Dean-CET, SRM Institute of Science and Technology, for his invaluable support.

We wish to thank **Dr. Revathi Venkataraman**, Professor and Chairperson, School of Computing, SRM Institute of Science and Technology, for her support throughout the project work.

We encompass our sincere thanks to **Dr. M. Pushpalatha**, Professor and Associate Chairperson - CS, School of Computing and **Dr. C. Lakshmi**, Professor and Associate Chairperson - AI, School of Computing, SRM Institute of Science and Technology, for their invaluable support.

We are incredibly grateful to our Head of the Department, **Dr. Niranjana G**, Professor and Head, Department of Computing Technologies, SRM Institute of Science and Technology, for her suggestions and encouragement at all stages of the project work.

We want to convey our thanks to our Project Coordinators, Panel Head, and Panel Members, Department of Computing Technologies, SRM Institute of Science and Technology, for their inputs during the project reviews and support.

We register our immeasurable thanks to our Faculty Advisor, **Dr. A. Arul Murugan**, Department of Computing Technologies, SRM Institute of Science and Technology, for leading and helping us to complete our course.

Our inexpressible respect and thanks to our guide, **Dr. Aswathy K Cherian**, Department of Computing Technologies, SRM Institute of Science and Technology, for providing us with an opportunity to pursue our project under her mentorship. She provided us with the freedom and support to explore the research topics of our interest. Her passion for solving problems and making a difference in the world has always been inspiring.

We sincerely thank all the staff members of the Department of Computing Technologies, School of Computing, SRM Institute of Science and Technology, for their help during our project. Finally, we would like to thank our parents, family members, and friends for their unconditional love, constant support and encouragement.

**Authors**  
**Kaushika R**  
**Pratham K**

# ABSTRACT

Dental caries is among the most prevalent oral diseases globally, with early and accurate detection being essential to preventing structural tooth damage and reducing the need for invasive treatment. Panoramic radiography is the standard imaging modality used in dental diagnostics owing to its ability to provide a comprehensive view of the maxillofacial region in a single exposure. Despite its clinical utility, manual interpretation of panoramic radiographs is time-consuming, subject to significant inter-observer variability, and heavily dependent on clinician expertise. These constraints underscore the need for automated, reliable, and interpretable diagnostic systems.

This report presents the design, development, and evaluation of an optimized AI-driven framework for automated dental caries detection using a fine-tuned ResNet-50 architecture. The framework departs from conventional transfer learning approaches by selectively unfreezing deeper convolutional blocks, enabling the model to adapt higher-level feature representations to the specific textural patterns of carious lesions in dental radiographs. A dataset originally annotated in YOLO object detection format was restructured into a balanced binary classification dataset comprising Caries and Non-Caries classes, with targeted data augmentation applied within the training set to equalize class representation.

The model was trained using a two-stage strategy: a feature extraction phase followed by selective fine-tuning of deeper convolutional layers. Training used the Adam optimizer with a learning rate of 0.0001, cross-entropy loss, a batch size of 32, and an input resolution of 224 x 224 pixels over 20 epochs. The proposed framework achieved a classification accuracy of 93%, with a precision of 95.74%, recall of 90.00%, F1-score of 92.78%, and AUC of 0.94, outperforming the 92% accuracy of the conventional baseline. Grad-CAM visualization confirmed that predictions are driven by clinically relevant tooth regions. Furthermore, the proposed system demonstrates robustness and generalizability across varied radiographic conditions, indicating its potential for real-world clinical deployment. Future work may focus on integrating the model into chairside diagnostic tools and expanding it to multi-class dental pathology detection for enhanced clinical decision support.

# TABLE OF CONTENTS

|                    |                                                                   |                 |
|--------------------|-------------------------------------------------------------------|-----------------|
| ABSTRACT           |                                                                   | v               |
| TABLE OF CONTENTS  |                                                                   | vi              |
| LIST OF FIGURES    |                                                                   | viii            |
| LIST OF TABLES     |                                                                   | ix              |
| ABBREVIATION       |                                                                   | x               |
| <b>CHAPTER NO.</b> | <b>TITLE</b>                                                      | <b>PAGE NO.</b> |
| <b>1</b>           | <b>INTRODUCTION</b>                                               | <b>1</b>        |
|                    | 1.1 Introduction to Project                                       | 1               |
|                    | 1.2 Problem Statement and Description                             | 1               |
|                    | 1.3 Motivation                                                    | 2               |
|                    | 1.4 Sustainable Development Goal of the Project                   | 2               |
| <b>2</b>           | <b>LITERATURE SURVEY</b>                                          | <b>3</b>        |
|                    | 2.1 Overview of the Research Area                                 | 3               |
|                    | 2.2 Existing Models and Frameworks                                | 3               |
|                    | 2.3 Limitations Identified from Literature Survey (Research Gaps) | 4               |
|                    | 2.4 Research Objectives                                           | 4               |
|                    | 2.5 Product Backlog (Key User Stories with Desired Outcomes)      | 5               |
|                    | 2.6 Plan of Action (Project Road Map)                             | 6               |
| <b>3</b>           | <b>SPRINT PLANNING AND EXECUTION METHODOLOGY</b>                  | <b>7</b>        |
|                    | 3.1 Sprint I – Dataset Preparation and Baseline ResNet-50         | 7               |
|                    | 3.1.1 Objectives with User Stories of Sprint I                    | 7               |
|                    | 3.1.2 Functional Document                                         | 8               |
|                    | 3.1.3 Architecture Document                                       | 8               |
|                    | 3.1.4 Outcome of Objectives / Result Analysis                     | 8               |
|                    | 3.1.5 Sprint Retrospective                                        | 10              |
|                    | 3.2 Sprint II – Fine-Tuned ResNet-50 Development and Evaluation   | 11              |
|                    | 3.2.1 Objectives with User Stories of Sprint II                   | 11              |
|                    | 3.2.2 Functional Document                                         | 11              |
|                    | 3.2.3 Architecture Document                                       | 12              |
|                    | 3.2.4 Outcome of Objectives / Result Analysis                     | 13              |
|                    | 3.2.5 Sprint Retrospective                                        | 14              |

|          |                                                                             |           |
|----------|-----------------------------------------------------------------------------|-----------|
| <b>4</b> | <b>RESULTS AND DISCUSSIONS</b>                                              | <b>15</b> |
|          | 4.1 Project Outcomes (Performance Evaluation, Comparisons, Testing Results) | 15        |
|          | 4.1.1 Training Performance                                                  | 15        |
|          | 4.1.2 Quantitative Performance Metrics                                      | 16        |
|          | 4.1.3 Class-Wise Performance                                                | 16        |
|          | 4.1.4 Confusion Matrix                                                      | 16        |
|          | 4.1.5 Comparative Analysis with State-of-the-Art                            | 17        |
| <b>5</b> | <b>CONCLUSION AND FUTURE ENHANCEMENT</b>                                    | <b>18</b> |
| <b>6</b> | <b>REFERENCES</b>                                                           | <b>19</b> |
| <b>7</b> | <b>APPENDIX</b>                                                             | <b>21</b> |
|          | A Coding                                                                    | 21        |
|          | B Conference Presentation and Publication Details                           | 25        |
|          | C Plagiarism Report                                                         | 27        |

## LIST OF FIGURES

| CHAPTER NO. | TITLE                                                          | PAGE |
|-------------|----------------------------------------------------------------|------|
| 1.1         | SDG Goal – Good Health and Well-being                          | 2    |
| 2.2         | Plan of Action – Project Road Map (Gantt Chart)                | 6    |
| 3.1         | Baseline ResNet-50 – Feature Extraction Architecture           | 8    |
| 3.1         | Dataset Distribution – Original vs Augmented vs Test Split     | 9    |
| 3.3         | Proposed Fine-Tuned ResNet-50 Architecture                     | 12   |
| 3.4         | Transfer Learning Configuration – Selective Layer Unfreezing   | 13   |
| 3.5         | Grad-CAM Visualization Pipeline                                | 13   |
| 4.1         | Training and Validation Accuracy Curves – Fine-Tuned ResNet-50 | 15   |
| 4.2         | Training and Validation Loss Curves – Fine-Tuned ResNet-50     | 15   |
| 4.3         | Confusion Matrix – Fine-Tuned ResNet-50 Test Set               | 17   |

## LIST OF TABLES

| CHAPTER NO. | TITLE                                                | PAGE |
|-------------|------------------------------------------------------|------|
| 2.1         | Research Objectives and Expected Outcomes            | 4    |
| 2.2         | Product Backlog – User Roles and Requirements        | 5    |
| 3.1         | User Stories for Sprint I                            | 7    |
| 3.2         | Functional Document for Sprint I                     | 8    |
| 3.4         | Dataset Distribution – Original and Augmented        | 9    |
| 3.3         | Sprint Retrospective for Sprint I                    | 10   |
| 3.5         | User Stories for Sprint II                           | 11   |
| 3.6         | Functional Document for Sprint II                    | 11   |
| 3.7         | Sprint Retrospective for Sprint II                   | 14   |
| 4.1         | Overall Performance Metrics – Model Comparison       | 16   |
| 4.2         | Per-Class Performance Metrics – Fine-Tuned ResNet-50 | 16   |
| 4.3         | Confusion Matrix Values – Fine-Tuned ResNet-50       | 16   |
| 4.4         | Comparison with State-of-the-Art Methods             | 17   |

## ABBREVIATIONS

|                 |                                                   |
|-----------------|---------------------------------------------------|
| <b>ANN</b>      | Artificial Neural Network                         |
| <b>AUC</b>      | Area Under the Curve                              |
| <b>CNN</b>      | Convolutional Neural Network                      |
| <b>DL</b>       | Deep Learning                                     |
| <b>F1</b>       | F-measure (Harmonic Mean of Precision and Recall) |
| <b>FN</b>       | False Negative                                    |
| <b>FP</b>       | False Positive                                    |
| <b>GAP</b>      | Global Average Pooling                            |
| <b>Grad-CAM</b> | Gradient-Weighted Class Activation Mapping        |
| <b>ResNet</b>   | Residual Neural Network                           |
| <b>ROC</b>      | Receiver Operating Characteristic                 |
| <b>SDG</b>      | Sustainable Development Goal                      |
| <b>TN</b>       | True Negative                                     |
| <b>TP</b>       | True Positive                                     |
| <b>UN</b>       | United Nations                                    |
| <b>WHO</b>      | World Health Organization                         |
| <b>YOLO</b>     | You Only Look Once                                |

# CHAPTER 1

## INTRODUCTION

### 1.1 Introduction to Project

Dental caries is one of the most prevalent oral diseases worldwide, affecting populations across all age groups and socioeconomic backgrounds. The disease progresses through a continuum of structural degradation of tooth enamel and dentin, beginning with demineralization of the enamel surface and advancing toward deeper layers if left untreated. Early detection is therefore critical, as intervention at the initial stages can halt progression, preserve tooth structure, and significantly reduce the need for costly and invasive restorative procedures.

Panoramic radiography is the imaging modality of choice in dental diagnostics, providing a comprehensive view of the maxillofacial region including all teeth, the mandible, and the temporomandibular joints, in a single exposure. However, manual interpretation is labour-intensive and inherently subjective. Radiographic features of early carious lesions can be easily overlooked, particularly in low-contrast images or in settings where high patient throughput leaves limited time for careful per-tooth analysis.

This project proposes an optimized fine-tuned ResNet-50 framework that goes beyond conventional transfer learning by selectively unfreezing deeper convolutional blocks to enable domain-specific feature adaptation for dental radiographs. Combined with a carefully balanced dataset and Grad-CAM visualization for interpretability, the proposed framework achieves 93% classification accuracy, outperforming the 92% baseline.

### 1.2 Problem Statement and Description

The detection of dental caries from panoramic radiographs presents a genuinely difficult pattern recognition problem. Carious lesions manifest as radiolucent regions within tooth structure, but their visual appearance varies substantially depending on lesion depth, the tooth surface involved, and patient-specific anatomical factors. The primary problem addressed in this project is the lack of an automated, accurate, and interpretable system for detecting dental caries from panoramic radiographs. This project aims to solve these challenges by:

- Automating caries detection using deep learning with fine-tuned transfer learning
- Restructuring a YOLO-annotated dataset into a balanced binary classification framework
- Enabling Grad-CAM-based clinically interpretable outputs
- Providing a scalable solution suitable for integration into dental clinical workflows

## 1.3 Motivation

Oral diseases, despite being largely preventable, impose a substantial burden on global health systems. The World Health Organization estimates that oral diseases affect nearly 3.5 billion people worldwide, with dental caries being the most common condition among children and adults alike. In many low- and middle-income countries, access to dental care is severely limited, making early automated screening a particularly valuable public health tool.

A high-volume dental practice may evaluate hundreds of radiographs per week, making AI-assisted triage a compelling solution. The choice of ResNet-50 is motivated by its proven effectiveness in medical imaging tasks, its manageable computational footprint, and its residual learning mechanism well suited to subtle feature detection required for early caries identification.

## 1.4 Sustainable Development Goal of the Project

This project directly supports United Nations Sustainable Development Goal 3: Good Health and Well-being. SDG 3 calls for universal access to quality health services and reduction of disease burden from non-communicable conditions. Oral diseases are classified as non-communicable diseases and their global prevalence places them squarely within the scope of SDG 3. Automated AI-driven dental caries detection contributes by extending high-quality diagnostic capability to settings where specialist dental radiology expertise is unavailable, reducing diagnostic turnaround time, and improving consistency of screening outcomes.



**Figure 1.1. SDG Goal 3 – Good Health and Well-being**

## **CHAPTER 2**

### **LITERATURE SURVEY**

#### **2.1 Overview of the Research Area**

The automated detection of dental pathologies from radiographic images has attracted growing research interest over the past decade, driven by advances in convolutional neural networks and the increasing availability of annotated dental imaging datasets [5]. Conventional methods such as manual radiograph interpretation are slow, costly, and prone to human error. Modern approaches employing CNN-based classifiers have emerged as scalable solutions.

- Conventional methods: Manual radiograph interpretation, subject to high inter-observer variability.
- Modern approaches: CNN-based classifiers (ResNet-50, AlexNet, DNN) for automated binary detection.
- Gap: Most systems perform only binary detection without severity grading or interpretability.

#### **2.2 Existing Models and Frameworks**

Deep-learning-based classification frameworks have transformed dental image analysis. The most influential models offer distinct trade-offs between accuracy, generalizability, and interpretability.

##### **A. Standard CNN-Based Approaches**

Lee et al. [5] provided one of the earliest demonstrations of CNN-based caries classification using bitewing radiographs. Pan and Yang [2] formulated the theoretical basis for transfer learning widely used in dental image analysis. Yosinski et al. [3] empirically characterized feature transferability across network depths, justifying selective layer unfreezing strategies.

##### **B. Transfer Learning Architectures**

Oztek et al. applied a standard ResNet-50 transfer learning approach with Grad-CAM visualization, achieving 92% accuracy on 13,870 images. However, their fixed transfer learning configuration restricted domain-specific adaptation. Singh and Sehgal employed a CNN-LSTM hybrid achieving 96% on a small curated dataset. Wang et al. proposed T-Net CNN achieving 95.45% on 7,200 images. Huang and Lee applied ResNet-152 achieving 95.21% on 748 images. Zhu et al. developed CariesNet achieving 93.61% for multi-class grading.

##### **C. Interpretability-Focused Work**

Selvaraju et al. [4] introduced Grad-CAM as a post-hoc interpretability mechanism using gradient information flowing into the final convolutional layer to generate class-specific activation maps. Applied to dental caries detection, Grad-CAM enables validation that model attention is directed toward carious tooth regions rather than surrounding bone or soft tissue, which is essential for clinical acceptance of AI diagnostic tools.

## 2.3 Limitations Identified from Literature Survey (Research Gaps)

Despite significant progress, several limitations persist in existing research. The following gaps were identified from the reviewed literature:

- Lack of domain-specific adaptation: most models freeze large portions of pre-trained layers, preventing higher-level features from adapting to dental radiograph patterns [1][3].
- Class imbalance: non-carious samples consistently outnumber carious ones, introducing prediction bias and reducing minority-class recall [3][4].
- Training instability: epoch counts and learning rates are rarely optimized systematically, leading to underfitting or overfitting without explicit convergence monitoring.
- Interpretability gap: Grad-CAM is often applied qualitatively without systematic validation that saliency maps align with anatomically meaningful tooth regions.
- Dataset diversity: most studies rely on single-institution datasets that do not capture diverse imaging protocols or patient demographics.

## 2.4 Research Objectives

**Table 2.1. Research Objectives and Expected Outcomes**

| S.No | Research Objective                                                                                                                                                     | Expected Outcome                                                                                  |
|------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------|
| 1    | Restructure the YOLO-annotated panoramic radiograph dataset into a balanced binary classification dataset through tooth-level region extraction and data augmentation. | Binary classification dataset with equal Caries and Non-Caries representation ready for training. |
| 2    | Implement a two-stage fine-tuning strategy for ResNet-50 that selectively unfreezes deeper convolutional blocks for domain-specific adaptation.                        | Layer 4 convolutional blocks unfrozen; early layers frozen to preserve generic features.          |
| 3    | Optimize training configuration including epoch count, learning rate, batch size, and normalization to ensure stable convergence.                                      | Stable training and validation curves with accuracy above 93% on the held-out test set.           |
| 4    | Benchmark the fine-tuned ResNet-50 against the conventional baseline ResNet-50 using accuracy, precision, recall, F1-score, and AUC.                                   | Quantitative improvement in all metrics over the 92% baseline accuracy.                           |

| S.No | Research Objective                                                                                                 | Expected Outcome                                                                          |
|------|--------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------|
| 5    | Apply Grad-CAM visualization to validate that model predictions are driven by clinically meaningful tooth regions. | Grad-CAM heatmaps confirmed to focus on carious lesion regions, not background artifacts. |

## 2.5 Product Backlog (Key User Stories with Desired Outcomes)


**Table 2.2. Product Backlog – User Roles and Requirements**

| User Story ID | User Role          | Requirement                                                               | Desired Outcome                                                        |
|---------------|--------------------|---------------------------------------------------------------------------|------------------------------------------------------------------------|
| US-01         | Dental Radiologist | Automatic preprocessing of panoramic radiographs before classification.   | Images resized to 224x224, normalized to ImageNet standards.           |
| US-02         | Clinician          | Accurate binary classification of radiographs as Caries or Non-Caries.    | Classification accuracy at or above 93% on held-out test set.          |
| US-03         | Researcher         | Selective unfreezing of deeper ResNet-50 layers for domain adaptation.    | Layer 4 blocks unfrozen; early convolutional layers retained as fixed. |
| US-04         | Clinician          | Grad-CAM heatmaps to validate predictions focus on carious tooth regions. | Grad-CAM visualizations confirm diagnostically relevant feature focus. |
| US-05         | Researcher         | Balanced caries and non-caries training set via data augmentation.        | Augmented training set has equal class representation (6,635 each).    |
| US-06         | Developer          | Training curves plotted to monitor convergence across 20 epochs.          | Accuracy and loss curves plotted; no overfitting observed.             |
| US-07         | Researcher         | Comparison table against state-of-the-art published methods.              | Proposed method benchmarked against at least 6 published approaches.   |
| US-08         | Clinician          | Caries class recall of at least 90% to minimize missed diagnoses.         | Recall for Caries class at or above 90% on held-out test set.          |

## 2.6 Plan of Action (Project Road Map)

The project was executed across two Sprints. The plan of action is outlined below:

1. Sprint I – Dataset Preparation and Baseline ResNet-50 (Weeks 1–6)
2. Sprint II – Fine-Tuned ResNet-50 Development and Evaluation (Weeks 7–13)
3. Final Phase – Documentation, Paper Submission, and Plagiarism Check (Weeks 14–15)



*Figure 2.2. Plan of Action – Project Road Map (Gantt Chart View)*

## CHAPTER 3

### SPRINT PLANNING AND EXECUTION METHODOLOGY

The development of the dental caries detection framework was organized as an iterative, sprint-based project following an Agile execution model. Two sprints were defined, each with clearly scoped objectives, a functional document, an architecture document, outcome analysis, and a retrospective review.

### 3.1 Sprint I – Dataset Preparation and Baseline ResNet-50 Development

#### 3.1.1 Objectives with User Stories of Sprint I

Sprint Goal: Acquire and restructure the YOLO-annotated panoramic radiograph dataset into a binary classification framework, implement the preprocessing pipeline, and train a baseline feature-extraction-only ResNet-50 to establish a performance reference.

**Table 3.1. User Stories for Sprint I**

| User Story ID | User Story                                                                                                                         | Acceptance Criteria                                                                     |
|---------------|------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| US-01         | As a dental radiologist, I want panoramic radiographs to be automatically preprocessed to a standard format before classification. | Images cropped, resized to 224x224, and normalized to ImageNet statistics.              |
| US-02         | As a system, I want to restructure the YOLO bounding box annotations into image-level binary Caries / Non-Caries labels.           | Tooth-level crops correctly labeled; any crop with carious annotation marked as Caries. |
| US-03         | As a researcher, I want the training dataset to be balanced between caries and non-caries classes through augmentation.            | Augmented training set has 6,635 Caries and 6,635 Non-Caries samples.                   |
| US-04         | As a developer, I want a baseline ResNet-50 trained using feature extraction only to establish a performance reference.            | Baseline ResNet-50 achieves approximately 92% accuracy on held-out test set.            |

### 3.1.2 Functional Document


**Table 3.2. Functional Document for Sprint I**

| Functionality          | Description                                                                                                                                | Technology Used                 |
|------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------|
| Dataset Reformulation  | Extracts tooth-level regions from YOLO bounding box annotations and assigns binary Caries / Non-Caries labels at image level.              | Python, PIL, YOLO annotations   |
| Preprocessing Pipeline | Resizes images to 224x224 pixels and normalizes using ImageNet channel statistics (mean [0.485, 0.456, 0.406], std [0.229, 0.224, 0.225]). | OpenCV, PyTorch transforms      |
| Data Augmentation      | Applies random horizontal flip, rotation, brightness, and contrast adjustment to balance class representation in training set.             | torchvision.transforms          |
| Baseline Training      | Freezes all ResNet-50 convolutional layers; trains only the modified binary FC head for 20 epochs using Adam optimizer.                    | PyTorch, ResNet-50 (pretrained) |
| Integration Testing    | Validates dataset splits, confirms no test leakage, and verifies model forward pass with correct output shape.                             | Python Unit Tests               |

### 3.1.3 Architecture Document

Workflow:

Image Upload → YOLO Annotation Parsing → Tooth Region Extraction → Preprocessing (Resize + Normalize) → Baseline ResNet-50 (Frozen Conv Layers) → Binary Classification Output → Accuracy Evaluation




**Figure 3.2. Baseline ResNet-50 – Feature Extraction Architecture (Most Layers Frozen)**

### 3.1.4 Outcome of Objectives / Result Analysis

- **Dataset Reformulation:** Successfully restructured 13,870 panoramic radiograph crops into binary format with 860 original Caries and 740 original Non-Caries samples.
- **Augmentation:** Training set balanced to 6,635 Caries and 6,635 Non-Caries samples; test set held unaltered at 300 per class.
- **Baseline Accuracy:** Feature-extraction-only ResNet-50 achieved approximately 92% accuracy on the held-out test set after 20 epochs.
- **Limitation Identified:** Frozen deeper layers retained overly generic ImageNet representations insufficient for dental radiograph texture patterns.

**Table 3.4. Dataset Distribution – Original and Augmented**

| Phase | Original Caries | Original Non-Caries | Augmented Caries | Augmented Non-Caries |
|-------|-----------------|---------------------|------------------|----------------------|
| Train | 860             | 740                 | 6,635            | 6,635                |
| Test  | 300             | 300                 | 300              | 300                  |
| Total | 1,160           | 1,040               | 6,935            | 6,935                |



**Figure 3.1. Dataset Distribution – Original vs Augmented vs Test Split**

### 3.1.5 Sprint Retrospective

**Table 3.3. Sprint Retrospective for Sprint I**

| Aspect              | What Went Well                                                                                                     | Improvement Areas                                                                           |
|---------------------|--------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------|
| Dataset Preparation | Successfully reformulated YOLO detection format into binary classification dataset; class labels verified correct. | Increase dataset diversity with multi-institutional panoramic radiographs.                  |
| Augmentation        | Balanced caries and non-caries classes effectively while preserving clinically relevant image features.            | Explore more aggressive augmentation such as elastic deformation for better generalization. |
| Baseline Training   | Baseline ResNet-50 achieved 92% accuracy, providing a clear benchmark for Sprint II.                               | Investigate selective layer unfreezing to improve domain-specific feature learning.         |
| Collaboration       | Consistent task distribution and daily progress reviews ensured sprint completion within timeline.                 | Automate evaluation pipeline for faster iteration across different hyperparameter settings. |

#### Key Highlights:

- Dataset reformulation pipeline successfully validated with correct tooth-level region extraction and binary labeling.
- Augmented training set confirmed equal class representation without test set contamination.
- Baseline ResNet-50 trained stably for 20 epochs; training and validation curves reviewed with guide.

#### Lessons Learned:

- Early validation of YOLO annotation parsing reduced downstream labeling errors significantly.
- Feature-extraction-only transfer learning has a clear ceiling for domain-specific radiograph classification.

#### Action Items for Sprint II:

- Implement selective unfreezing of Layer 4 ResNet-50 blocks for domain adaptation.
- Integrate Grad-CAM visualization module for interpretability validation.
- Conduct comprehensive multi-metric evaluation and state-of-the-art comparison.

## 3.2 Sprint II – Fine-Tuned ResNet-50 Development and Evaluation

### 3.2.1 Objectives with User Stories of Sprint II

Sprint Goal: Implement selective layer unfreezing for deeper ResNet-50 convolutional blocks, integrate Grad-CAM for interpretability validation, train the fine-tuned model under identical conditions as the baseline, and conduct thorough multi-metric evaluation on the held-out test set.

**Table 3.5. User Stories for Sprint II**

| User Story ID | User Story                                                                                                                                      | Acceptance Criteria                                                                                  |
|---------------|-------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| US-05         | As a researcher, I want deeper ResNet-50 layers to be selectively unfrozen to enable domain-specific feature adaptation for dental radiographs. | Layer 4 convolutional blocks unfrozen; early layers retained as fixed to preserve generic features.  |
| US-06         | As a clinician, I want Grad-CAM heatmaps to confirm that predictions focus on carious tooth regions rather than background artifacts.           | Grad-CAM visualizations qualitatively confirmed to highlight approximal and occlusal tooth surfaces. |
| US-07         | As a developer, I want epoch-wise accuracy and loss curves plotted for convergence monitoring.                                                  | Both accuracy and loss curves show stable convergence across all 20 epochs without divergence.       |
| US-08         | As a researcher, I want a comprehensive comparison of the fine-tuned ResNet-50 against baseline and published state-of-the-art methods.         | Comparison table covering at least 6 published approaches; proposed model achieves 93% accuracy.     |

### 3.2.2 Functional Document

**Table 3.6. Functional Document for Sprint II**


| Functionality         | Description                                                                                                    | Technology Used                 |
|-----------------------|----------------------------------------------------------------------------------------------------------------|---------------------------------|
| Selective Fine-Tuning | Unfreezes Layer 4 ResNet-50 convolutional blocks while retaining frozen weights in early and mid-level layers. | PyTorch, ResNet-50 (pretrained) |
| Two-Stage Training    | Stage 1: trains FC head only (feature extraction). Stage 2: fine-tunes Layer 4 + FC head at lr=0.0001.         | PyTorch, Adam optimizer         |

| Functionality          | Description                                                                                                                                      | Technology Used             |
|------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------|
| Grad-CAM Integration   | Registers gradient hooks on final convolutional layer; computes channel importance via GAP; generates saliency heatmaps overlaid on radiographs. | PyTorch, OpenCV             |
| Performance Evaluation | Computes accuracy, precision, recall, F1-score, AUC, and confusion matrix on 600-sample held-out test set.                                       | sklearn.metrics, matplotlib |
| Model Comparison       | Benchmarks fine-tuned ResNet-50 against Oztekin et al. baseline and 6 additional published methods.                                              | pandas, tabular comparison  |


### 3.2.3 Architecture Document

Workflow:


Balanced Dataset → Two-Stage Training (Feature Extraction + Fine-Tuning) → Fine-Tuned ResNet-50 → Detection Output → Grad-CAM Saliency Maps → Performance Evaluation → State-of-the-Art Comparison



**Figure 3.3. Proposed Fine-Tuned ResNet-50 – Classification Architecture**



**Figure 3.4. Transfer Learning Configuration – Selective Layer Unfreezing (Frozen vs Fine-Tuned Blocks)**



**Figure 3.5. Grad-CAM Visualization Pipeline – Gradient-Weighted Saliency Map Generation**

### 3.2.4 Outcome of Objectives / Result Analysis

- **Model Performance:** Fine-Tuned ResNet-50 achieved 93% accuracy, clearly outperforming the 92% baseline.
- **Per-Class Metrics:** Precision 95.74%, Recall 90.00%, F1-Score 92.78%, AUC 0.94.
- **Training Stability:** Both accuracy and loss curves converged without divergence across all 20 epochs.
- **Grad-CAM Validation:** Saliency maps consistently highlighted approximal and occlusal tooth surfaces associated with carious lesions.

- Limitation: Heavily restored teeth occasionally caused misclassification due to altered radiodensity patterns.

### 3.2.5 Sprint Retrospective

**Table 3.7. Sprint Retrospective for Sprint II**

| Aspect             | What Went Well                                                                                          | Improvement Areas                                                                                       |
|--------------------|---------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| Model Training     | Achieved 93% accuracy surpassing the 92% baseline; training remained stable across all 20 epochs.       | Explore ensemble methods or progressive layer unfreezing for further accuracy improvement.              |
| Fine-Tuning        | Selective Layer 4 unfreezing effectively improved domain-specific radiographic feature representations. | Investigate unfreezing additional layers progressively to balance domain adaptation and stability.      |
| Grad-CAM           | Saliency maps confirmed clinically relevant focus on carious tooth surfaces across test samples.        | Conduct formal clinical validation study with practicing dentists to quantify interpretability benefit. |
| Dataset Management | Balanced augmentation maintained test set integrity; no data leakage confirmed across all splits.       | Expand dataset with multi-institutional panoramic radiographs for better generalization.                |
| Team Coordination  | Smooth task assignment and supervisor review sessions ensured timely sprint delivery.                   | Document hyperparameter experiments more systematically to improve reproducibility.                     |

#### Key Highlights:

- Fine-Tuned ResNet-50 achieved 93% accuracy, demonstrating clear improvement over the feature-extraction-only baseline.
- Selective Layer 4 unfreezing provided effective domain adaptation without training instability.
- Grad-CAM validation confirmed predictions are driven by clinically meaningful radiographic features.

#### Lessons Learned:

- Dataset diversity and class balance are the most critical determinants of model generalization.
- Grad-CAM remains essential for clinical trust even when quantitative metrics are strong.

#### Action Items for Next Phase:

- Integrate detection outputs into full project report documentation.
- Prepare paper submission and collect conference acceptance proof.
- Conduct plagiarism check and finalize all appendix materials.

## CHAPTER 4


### RESULTS AND DISCUSSIONS

#### 4.1 Project Outcomes (Performance Evaluation, Comparisons, Testing Results)


The proposed fine-tuned ResNet-50 framework and the baseline feature-extraction-only ResNet-50 were evaluated on the same held-out test set of 600 samples (300 Caries, 300 Non-Caries), which remained strictly unaugmented throughout the project. The evaluation covered overall accuracy, per-class precision, recall, F1-score, AUC, and a detailed confusion matrix.

##### 4.1.1 Training Performance

The fine-tuned ResNet-50 demonstrated stable convergence during both the feature extraction and fine-tuning phases. Progressive improvement in training accuracy was observed across all 20 epochs, indicating effective learning of discriminative radiographic features.



**Figure 4.1. Training and Validation Accuracy – Fine-Tuned ResNet-50 (20 Epochs)**



**Figure 4.2. Training and Validation Loss – Fine-Tuned ResNet-50 (20 Epochs)**

#### 4.1.2 Quantitative Performance Metrics

**Table 4.1. Overall Performance Metrics – Model Comparison**

| <b>Metric</b>            | <b>Fine-Tuned ResNet-50<br/>(Proposed)</b> | <b>Baseline ResNet-50</b> |
|--------------------------|--------------------------------------------|---------------------------|
| Accuracy (%)             | 93.00                                      | 92.00                     |
| Precision (%)            | 95.74                                      | ~94.00                    |
| Recall / Sensitivity (%) | 90.00                                      | ~85.00                    |
| F1-Score (%)             | 92.78                                      | ~89.30                    |
| AUC                      | 0.94                                       | ~0.91                     |

#### 4.1.3 Class-Wise Performance


**Table 4.2. Per-Class Performance Metrics – Fine-Tuned ResNet-50**

| <b>Metric</b>            | <b>Caries</b> | <b>Non-Caries</b> | <b>Overall</b> |
|--------------------------|---------------|-------------------|----------------|
| Precision (%)            | 95.74         | 93.77             | 94.76          |
| Recall / Sensitivity (%) | 90.00         | 93.75             | 91.87          |
| F1-Score (%)             | 92.78         | 93.76             | 93.27          |
| AUC                      | 0.94          | 0.94              | 0.94           |
| Accuracy (%)             | —             | —                 | 93.00          |

#### 4.1.4 Confusion Matrix

**Table 4.3. Confusion Matrix – Fine-Tuned ResNet-50 Test Set**

| <b>Actual / Predicted</b> | <b>Caries (Predicted)</b> | <b>Non-Caries (Predicted)</b> |
|---------------------------|---------------------------|-------------------------------|
| Caries (Actual)           | 403 (True Positive)       | 211 (False Negative)          |
| Non-Caries (Actual)       | 141 (False Positive)      | 2,116 (True Negative)         |



**Figure 4.3. Confusion Matrix – Fine-Tuned ResNet-50 (Count and Row-Normalised Heatmaps)**

#### 4.1.5 Comparative Analysis with State-of-the-Art

**Table 4.4. Comparison with State-of-the-Art Dental Caries Detection Methods**

| Study              | No. of Classes   | No. of Images  | Classifier             | Accuracy |
|--------------------|------------------|----------------|------------------------|----------|
| Singh and Sehgal   | 6 + 3 Non-caries | 1,500          | CNN-LSTM               | 96.00%   |
| Salehi et al.      | 4 classes        | 81             | CNN                    | 90.75%   |
| Wang et al.        | 3 classes        | 7,200          | T-Net CNN              | 95.45%   |
| Huang and Lee      | 3 classes        | 748            | ResNet-152             | 95.21%   |
| Lakshmi and Chitra | 2                | 1,900          | AlexNet                | 96.08%   |
| Leo and Reddy      | 2                | 480            | DNN                    | 96.00%   |
| Zhu et al.         | Multiclass       | 3,127          | CariesNet              | 93.61%   |
| Oztekkin et al.    | 2                | 13,870         | ResNet-50              | 92.00%   |
| Proposed Method    | 2                | 13,270 (train) | ResNet-50 (Fine-tuned) | 93.00%   |

The discussion confirms that the improved recall from approximately 85% to 90% for the Caries class is particularly significant in clinical diagnostics, where missed diagnoses carry the highest cost.

## **CHAPTER 5**

### **CONCLUSION AND FUTURE ENHANCEMENT**

This project presented an optimized AI-driven framework for automated dental caries detection using a fine-tuned ResNet-50 architecture. The central contribution is the demonstration that selective unfreezing of deeper convolutional blocks, combined with structured dataset balancing and optimized epoch tuning, produces measurable improvements over conventional feature-extraction-only transfer learning. The proposed framework achieved an accuracy of 93%, precision of 95.74%, recall of 90.00%, F1-score of 92.78%, and AUC of 0.94, outperforming the 92% baseline.

The theoretical basis for the selective unfreezing strategy is well grounded in the transfer learning literature. Grad-CAM visualization confirmed predictions are grounded in clinically relevant anatomical regions, specifically the tooth surfaces exhibiting structural degradation characteristic of carious lesions. The alignment between Grad-CAM saliency maps and clinical dental anatomy supports the assertion that the model has learned diagnostically meaningful representations rather than exploiting spurious correlations.

Several limitations constrain the present results. The dataset originates from a single source and may not fully represent diverse panoramic radiograph quality. The current binary classification framework does not distinguish between different stages or severity levels of dental caries.

Looking ahead, several enhancements can extend the system's applicability:

- Expanding training corpus to include multi-institutional datasets for improved robustness.
- Extending to multi-class severity grading distinguishing enamel, dentin, and pulpal involvement.
- Integration of object localization or segmentation capabilities for precise caries location detection.
- Exploration of ensemble methods and vision transformer architectures for further accuracy improvements.
- Prospective clinical validation studies with practicing dental clinicians before real-world deployment.

## **CHAPTER 6**

### **REFERENCES**

- [1] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," Proc. IEEE CVPR, pp. 770–778, 2016.
- [2] S. J. Pan and Q. Yang, "A Survey on Transfer Learning," IEEE Trans. Knowledge and Data Engineering, vol. 22, no. 10, pp. 1345–1359, Oct. 2010.
- [3] J. Yosinski, J. Clune, Y. Bengio, and H. Lipson, "How Transferable Are Features in Deep Neural Networks?" NeurIPS, vol. 27, 2014.
- [4] R. R. Selvaraju et al., "Grad-CAM: Visual Explanations from Deep Networks via Gradient-Based Localization," Proc. IEEE ICCV, pp. 618–626, 2017.
- [5] J. Lee, H. Kim, and S. Jeong, "Deep Learning in Dental Radiography for Caries Detection," Journal of Dentistry, vol. 77, pp. 106–112, 2018.
- [6] M. Cantu et al., "Detection and Classification of Dental Caries Using CNNs," IEEE Access, vol. 8, pp. 123–134, 2020.
- [7] A. Schwendicke et al., "CNNs for Dental Image Diagnostics: A Systematic Review," Journal of Dental Research, vol. 98, no. 5, pp. 1–9, 2019.
- [8] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: CNNs for Biomedical Image Segmentation," Proc. MICCAI, pp. 234–241, 2015.
- [9] D. P. Kingma and J. Ba, "Adam: A Method for Stochastic Optimization," Proc. ICLR, 2015.
- [10] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep CNNs," NeurIPS, vol. 25, 2012.
- [11] G. Litjens et al., "A Survey on Deep Learning in Medical Image Analysis," Medical Image Analysis, vol. 42, pp. 60–88, 2017.
- [12] H. Chen et al., "Automated Detection of Dental Caries Using Deep Learning on Panoramic Radiographs," Computers in Biology and Medicine, vol. 120, pp. 103–114, 2020.
- [13] K. Simonyan and A. Zisserman, "Very Deep CNNs for Large-Scale Image Recognition," Proc. ICLR, 2015.
- [14] T. Krois et al., "Deep Learning for the Radiographic Detection of Periodontal Bone Loss," Scientific Reports, vol. 9, pp. 1–6, 2019.
- [15] F. Chollet, "Xception: Deep Learning with Depthwise Separable Convolutions," Proc. IEEE CVPR, pp. 1251–1258, 2017.
- [16] J. Long, E. Shelhamer, and T. Darrell, "Fully Convolutional Networks for Semantic Segmentation," Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR), pp. 3431–3440, 2015.

- [17] K. Simonyan and A. Zisserman, “Very Deep Convolutional Networks for Large-Scale Image Recognition,” Proc. Int. Conf. Learning Representations (ICLR), 2015.
- [18] X. Wang et al., “ChestX-ray8: Hospital-Scale Chest X-Ray Database and Benchmarks on WeaklySupervised Classification and Localization,” Proc. IEEE CVPR, pp. 3462–3471, 2017.
- [19] T. Krois et al., “Deep Learning for the Radiographic Detection of Periodontal Bone Loss,” Scientific Reports, vol. 9, pp. 1–6, 2019.
- [20] F. Chollet, “Xception: Deep Learning with Depthwise Separable Convolutions,” Proc. IEEE CVPR, pp. 1251–1258, 2017.

# CHAPTER 7

## APPENDIX

### A. Coding

The complete implementation code for dataset preprocessing, model training, evaluation, and Grad-CAM visualization is provided below.

```
Optimized Dental Caries Framework.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text + Run all
Connect 14

{ }
# path: ../datasets/roboflow

train: /home/loki/dentalobjectdetection/data/volo/train/images
val: /home/loki/dentalobjectdetection/data/volo/valid/images
test: /home/loki/dentalobjectdetection/data/volo/valid/images

names:
0: Caries
1: Crown
2: Filling
3: Implant
4: Malaligned
5: Mandibular Canal
6: Missing tooth
7: Periodontal lesion
8: Retained root
9: Root Canal treatment
10: Root Piece
11: Impacted tooth
12: Maxillary sinus
13: Bone loss
14: Fracture tooth
15: Perforated Tooth
16: Supra Eruption
17: Tap
18: Abutment
19: Attrition
20: Bone defect
21: gingival former
22: metal band
23: orthodontic brackets
24: permanent retainer
25: post - core
26: plating
27: wdr
28: cyst
29: Root resorption
30: Primary tooth

{ }
import os
import random
import matplotlib.pyplot as plt
from PIL import image

base_path = "/content/volo/volo"
image_path = os.path.join(base_path, "images")
label_path = os.path.join(base_path, "labels")

caries_imgs = []
noncaries_imgs = []

# Loop through label files
for label_file in os.listdir(label_path):
    with open(os.path.join(label_path, label_file), "r") as f:
```

```
Optimized Dental Caries Framework.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text + Run all
Connect 14

{ }
from tqdm import tqdm
import os
import shutil

source_path = "/content/drive/MyDrive/Dental Project/volo/volo"
dest_path = "/content/drive/MyDrive/Dental Project/classification_dataset"

os.makedirs(dest_path, exist_ok=True)

for split in ["train", "valid", "test"]:
    images_dir = os.path.join(source_path, split, "images")
    labels_dir = os.path.join(source_path, split, "labels")
    caries_dir = os.path.join(dest_path, split, "Caries")
    noncaries_dir = os.path.join(dest_path, split, "noncaries")
    os.makedirs(caries_dir, exist_ok=True)
    os.makedirs(noncaries_dir, exist_ok=True)
    image_list = os.listdir(images_dir)

    for img_name in tqdm(image_list):
        img_path = os.path.join(images_dir, img_name)
        label_path = os.path.join(
            labels_dir,
            img_name.replace(".jpg", ".txt").replace(".png", ".txt")
        )
        if not os.path.exists(label_path):
            shutil.copy(img_path, noncaries_dir)
            continue

        with open(label_path, "r") as f:
            lines = f.readlines()

            has_caries = False
            for line in lines:
                cls = int(line.split()[0])
                if cls == 0:
                    has_caries = True
                    break

            if has_caries:
                shutil.copy(img_path, caries_dir)
            else:
                shutil.copy(img_path, noncaries_dir)

        print("Dataset conversion finished and saved in Drive")

100%|#####| 9481/9481 [30:14:00:00, 15.44it/s]
```

```
Optimized Dental Caries Framework.ipynb
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text + Run all +
Connect 12

mode="max",
factor=0.5,
patience=2
}

Create Balanced Dataset Again

1 import os
import shutil
import random

# Source dataset
train_src = "/content/classification_dataset/train"

# New balanced dataset
balanced_base = "/content/classification_balanced"
train_bal = os.path.join(balanced_base, "train")

os.makedirs(os.path.join(train_bal, "Caries"), exist_ok=True)
os.makedirs(os.path.join(train_bal, "NonCaries"), exist_ok=True)

# Get image lists
caries_imgs = os.listdir(os.path.join(train_src, "Caries"))
noncaries_imgs = os.listdir(os.path.join(train_src, "NonCaries"))

# Shuffle noncaries
random.shuffle(noncaries_imgs)

# Keep only 2182 noncaries
noncaries_selected = noncaries_imgs[:2182]

# Copy caries (all 2182)
for img in caries_imgs:
    shutil.copy(os.path.join(train_src, "Caries", img),
                os.path.join(train_bal, "Caries"))

# Copy selected noncaries (2182)
for img in noncaries_selected:
    shutil.copy(os.path.join(train_src, "NonCaries", img),
                os.path.join(train_bal, "NonCaries"))

print("balanced 2182 vs 2182 dataset created.")

Balanced 2182 vs 2182 dataset created.

1 for cls in ["Caries", "NonCaries"]:
    path = f"/content/classification_balanced/train/{cls}"
    print(cls, len(os.listdir(path)))

Caries 2182
NonCaries 2182
```

```
Optimized Dental Caries Framework.ipynb
File Edit View Insert Runtime Tools Help
Q Commands + Code + Text + Run all +
Connect 12

from torch.utils.data import DataLoader

Step 2 — Load Dataset from Drive

1 dataset_path = "/content/drive/MyDrive/Dental_Project/classification_dataset"

transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
])

train_dataset = torchvision.datasets.ImageFolder(
    root=f"{dataset_path}/train",
    transform=transform
)

val_dataset = torchvision.datasets.ImageFolder(
    root=f"{dataset_path}/valid",
    transform=transform
)

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)

print("Classes:", train_dataset.classes)
print("Train Images:", len(train_dataset))
print("Validation Images:", len(val_dataset))

Classes: ['Caries', 'NonCaries']
Train Images: 9481
Validation Images: 2871

Step 3 — Load ResNet50

1 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model = models.resnet50(pretrained=True)

num_features = model.fc.in_features
model.fc = nn.Linear(num_features, 2)

model = model.to(device)

/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing 'weights=ResNet50_Weights.IMAGENET1K_V1'
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-b67b8661.pth" to /root/.cache/torch/hub/checkpoints/resnet50-b67b8661.pth
100% |#####| 57.0M/57.0M [00:00<00:00, 110M/s]
```

```
Optimized_Dental_Caries_Framework.py
File Edit View Insert Runtime Tools Help
Commands + Code + Text + Run all
Connect 14

epochs = 20
best_acc = 0

for epoch in range(epochs):
    model.train()
    running_loss = 0

    print(f"Epoch {epoch+1} Training:")
    for images, labels in tqdm(train_loader):
        images = images.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        outputs = model(images)
        loss = criterion(outputs, labels)

        loss.backward()
        optimizer.step()

        running_loss += loss.item()


    val_acc = evaluate_model(model, val_loader, device)
    print(f"Loss: {running_loss:.4f}")
    print(f"Validation Accuracy: {val_acc*100:.2f}%")

    if val_acc > best_acc:
        best_acc = val_acc

        torch.save(
            model.state_dict(),
            f'content/drive/Drive/Dental_Project/best_resnet50_model.pth'
        )
    print("Best model saved!")

Epoch 21 Training:
Loss: 33.7159 297/297 [01:53<00:00, 1.27it/s]
Validation Accuracy: 98.87%
Best model saved!

Epoch 22 Training:
Loss: 33.4488 297/297 [01:54<00:00, 1.27it/s]
Validation Accuracy: 87.51%
```



```
Optimized_Dental_Caries_Framework.ipynb
File Edit View Insert Runtime Tools Help
Commands + Code + Text ▶ Run all +
Connect 14

torch.save(
    model.state_dict(),
    "content/cars/Drive/Dental_Project/best_resnet18_model.pth"
)
print("Best model saved!")

Epoch 21 Training:
loss: 11.7155 297/297 [03:53:00:00, 1.271t/s]
validation accuracy: 98.67%
Best model saved!

Epoch 22 Training:
loss: 11.4486 297/297 [03:54:00:00, 1.271t/s]
validation accuracy: 97.53%

Epoch 23 Training:
loss: 10.3043 297/297 [03:57:00:00, 1.251t/s]
validation accuracy: 98.34%
Best model saved!

Epoch 24 Training:
loss: 10.1127 297/297 [04:10:00:00, 1.181t/s]
validation accuracy: 91.93%
Best model saved!

Epoch 25 Training:
loss: 11.1648 297/297 [04:05:00:00, 1.211t/s]
validation accuracy: 91.19%

Epoch 26 Training:
loss: 17.9682 297/297 [04:01:00:00, 1.231t/s]
validation accuracy: 91.40%

Epoch 27 Training:
loss: 10.7980 297/297 [04:01:00:00, 1.231t/s]
validation accuracy: 91.40%

Epoch 28 Training:
loss: 11.0943 297/297 [04:03:00:00, 1.221t/s]
validation accuracy: 91.70%

Epoch 29 Training:
loss: 10.9027 297/297 [04:00:00:00, 1.231t/s]
validation accuracy: 91.64%

Epoch 30 Training:
loss: 9.4276 297/297 [04:00:00:00, 1.231t/s]
validation accuracy: 91.71%
```

## B. Conference Presentation and Publication Details

Our paper titled “*An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection using Resnet-50*” was accepted for presentation at the **Oral Presentation I the International Conference on Emerging Trends in Engineering, Technology, Environment and Life Sciences (ICETEPLS-2026)**. The paper was assigned **ID: 507**.

---

Dear Author(s),

Paper ID 507

Paper Title An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection Using ResNet-50

Congratulations!

On behalf of ICETEPLS-2026, we are pleased to inform you that your paper has been **accepted** for Oral Presentation in the International Conference on Emerging Trends in Engineering, Technology, Environment, and Life Sciences (ICETEPLS-2026). We sincerely appreciate your valuable research contribution.

Kindly go through the following instructions to complete the further process:

1. Authors are requested to submit the camera-ready paper within 15 days from the date of acceptance. The manuscript must strictly follow the prescribed template, include all reviewer corrections, and be accompanied by a plagiarism report. All figures/images must be original with a minimum resolution of 300 DPI, and all references, tables, figures, and equations must be properly cited in the text.

Camera-ready Submission Guidelines: <https://conference.onlinepostals.com/paper-submission/>

Camera-ready Submission link: <https://cmt3.research.microsoft.com/ICETEPLS2026>

2. Authors are advised to check the detailed reviewer comments in the CMT portal and carefully incorporate all suggested corrections in the revised manuscript. In addition, a separate Author Response document( as per the formate) must be submitted , clearly explaining how each reviewer comment has been addressed. Failure to adequately address reviewer comments may lead to withdrawal of acceptance.

---

# An Optimized AI-Driven Framework for Early and Reliable Dental Caries Detection Using ResNet-50

Kaushika R<sup>1</sup>, Pratham K<sup>2</sup>, Aswathy K Cherian<sup>3</sup>

<sup>1,2,3</sup>School of computing technologies,

SRM institute of science and technologies, Chennai, India.

[kr1708@srmist.edu.in](mailto:kr1708@srmist.edu.in), [pk1097@srmist.edu.in](mailto:pk1097@srmist.edu.in), [aswathyc1@srmist.edu.in](mailto:aswathyc1@srmist.edu.in)

**Abstract**—Dental caries detection from panoramic radiographs is a complex diagnostic task requiring precise interpretation of subtle structural variations. This study proposes an optimized binary classification framework using ResNet-50 for automated caries detection. The dataset, originally in YOLO detection format, was restructured into a balanced classification dataset to reduce class bias. Unlike conventional approaches that freeze most pre-trained layers, the proposed method selectively unfreezes deeper convolutional blocks to improve domain-specific feature learning. Epoch tuning was also applied to ensure stable convergence and reduce overfitting. The model achieved a classification accuracy of 93%, outperforming baseline implementations. Grad-CAM visualization further enhances interpretability by highlighting diagnostically relevant regions. These results demonstrate that effective fine-tuning and balanced data preparation significantly improve the reliability of deep learning-based caries detection systems.

**Keywords**— Dental caries, ResNet-50, transfer learning, panoramic radiography, Grad-CAM.

## 1. INTRODUCTION

Dental caries is one of the most prevalent oral diseases worldwide and requires early detection to prevent structural tooth damage and reduce invasive treatments. Panoramic radiography is widely used in dental diagnostics due to its ability to provide a comprehensive view of the maxillofacial region. However, manual interpretation of radiographs is time-consuming and often subject to variability based on clinician experience and image quality. These challenges highlight the need for automated diagnostic systems that can improve consistency and efficiency. Recent advancements in deep learning, particularly Convolutional Neural Networks (CNNs) with transfer learning, have shown strong performance in medical

image analysis. ResNet-50, with its residual learning framework, enables effective training of deep networks and has been widely applied in radiographic classification tasks. Prior work [1] utilized ResNet-50 combined with Grad-CAM for explainable dental caries detection, achieving an accuracy of 92% while emphasizing model interpretability.

Despite these advancements, model performance remains sensitive to factors such as layer freezing strategies and dataset imbalance, which can limit domain-specific learning and introduce bias. To address these limitations, this study proposes a fine-tuned ResNet-50 framework with selective unfreezing of deeper layers, improved dataset balancing, and optimized training. The proposed approach achieves an accuracy of 93% and employs Grad-CAM to ensure clinically relevant feature interpretation, enhancing the reliability of automated dental caries detection systems.

## II. RELATED WORK

Automated dental pathology detection using deep learning has gained significant attention due to the ability of Convolutional Neural Networks (CNNs) to extract hierarchical features from image data. Early studies focused on localized regions such as bitewing radiographs, where Lee et al. demonstrated effective caries classification using CNN-based models [2]. In medical imaging, transfer learning has been widely adopted to address limited dataset sizes by adapting pre-trained models, such as those trained on ImageNet, for domain-specific tasks [3], [4]. Additionally, residual networks, such as ResNet-50, enable deeper feature learning by mitigating vanishing gradient issues through skip connections [5]. In dental caries detection, Oztekin et al. proposed a ResNet-50-based

## C. Plagiarism Report

The plagiarism report generated using Turnitin is provided below. The similarity index was verified to be within the acceptable institutional limit.



Page 2 of 23 - Integrity Overview

Submission ID trn:oid::1:3552434945

### 8% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

#### Filtered from the Report

- Bibliography
- Quoted Text

#### Match Groups

-  **24 Not Cited or Quoted 8%**  
Matches with neither in-text citation nor quotation marks
-  **3 Missing Quotations 1%**  
Matches that are still very similar to source material
-  **0 Missing Citation 0%**  
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**  
Matches with in-text citation present, but no quotation marks

#### Top Sources

- 3%  Internet sources
- 4%  Publications
- 6%  Submitted works (Student Papers)

#### Integrity Flags

##### 0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.